

**Surakshith Shetty-53026240013**

**Practical 5 : To analyze the impact of marketing mix variables on sales and develop a predictive time series model using SARIMAX for improved sales forecasting**

```

    \\"min\": 37.65617447,\n          \\"max\": 240.291967,\n    \\"num_unique_values\": 100,\n          \\"samples\": [\n      109.4686053,\n          130.5990849,\n          167.5356191\n      ],\n          \\"semantic_type\": \"\","n\n    \\"description\": \"\n      }\n    }\n  ]\n}\n},\"type\":\"dataframe\",\"variable_name\":\"data\"}

from statsmodels.formula.api import ols

model = ols('Sales~Base_Price+TV+Radio+InStore', data=data).fit()

model.summary()

<class 'statsmodels.iolib.summary.Summary'>
"""
                OLS Regression Results
=====

Dep. Variable:                 Sales   R-squared:
0.648
Model:                          OLS   Adj. R-squared:
0.633
Method:                         Least Squares   F-statistic:
43.65
Date:             Mon, 10 Mar 2025   Prob (F-statistic):
9.64e-21
Time:                  09:54:03   Log-Likelihood:
-825.76
No. Observations:            100   AIC:
1662.
Df Residuals:                  95   BIC:
1675.
Df Model:                      4

Covariance Type:            nonrobust
=====

            coef    std err          t      P>|t|      [0.025
0.975]
-----
Intercept  4.757e+04    3019.360     15.756      0.000      4.16e+04
5.36e+04
Base_Price -1952.5944    191.249    -10.210      0.000     -2332.271
-1572.918
TV           7.4444     2.217      3.357      0.001      3.042
11.847
Radio         1.1928     1.109      1.076      0.285     -1.008

```

```

3.394
InStore      35.0531     7.323      4.787      0.000     20.515
49.591
=====
=====
Omnibus:          3.728  Durbin-Watson:
0.898
Prob(Omnibus):    0.155  Jarque-Bera (JB):
3.098
Skew:             -0.332  Prob(JB):
0.212
Kurtosis:         3.551   Cond. No.
9.63e+03
=====
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 9.63e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```

"""
exog_data = data[['Base_Price', 'TV', 'Radio', 'InStore']]
scaler = StandardScaler()
scaled_exog_data = scaler.fit_transform(exog_data)
scaled_exog_data = pd.DataFrame(scaled_exog_data,
columns=exog_data.columns, index=exog_data.index)
scaled_exog_data.head()

{
  "summary": {
    "name": "scaled_exog_data",
    "rows": 100,
    "fields": [
      {
        "column": "Base_Price",
        "properties": {
          "dtype": "number",
          "std": 1.0050378152592123,
          "min": -2.970392187747817,
          "max": 1.9135544692249764,
          "num_unique_values": 17,
          "samples": [
            -0.48830835385693155,
            1.340610813230131,
            0.9902662253116309
          ],
          "semantic_type": "\",
          "description": "\n"
        }
      },
      {
        "column": "TV",
        "properties": {
          "dtype": "number",
          "std": 1.005037815259212,
          "min": -2.368876090359765,
          "max": 2.3093569987797444,
          "num_unique_values": 100,
          "samples": [
            -0.7109493937683404,
            -0.22311203976548163,
            0.6296381794091025
          ],
          "semantic_type": "\",
          "description": "\n"
        }
      },
      {
        "column": "Radio",
        "properties": {
          "dtype": "number",
          "std": 1.005037815259212,
          "min": -2.9655049218752,
          "max": 1.6440882209359917,
          "num_unique_values": 53,
          "samples": [
            0.8238348045460303
          ]
        }
      }
    ]
  }
}
```

```

0.11911003835183809,\n                 -0.6433790529402388\n           ],\n  \"semantic_type\": \"\",\\n      \"description\": \"\"\n},\\n  {\n    \"column\": \"InStore\",\\n      \"properties\":\n      {\n        \"dtype\": \"number\",\\n          \"std\":\n1.005037815259212,\n        \"min\": -1.595425936104774,\n        \"max\": 2.6195319934669197,\n        \"num_unique_values\": 100,\n        \"samples\": [\n          -1.0593763830685705,\n          -0.9700592948981115,\n          0.3501428338700144\n        ],\n        \"semantic_type\": \"\",\\n          \"description\": \"\"\n      }\n    }\n  ]\\n},\"type\":\"dataframe\",\"variable_name\":\"scaled_exog_data\"}

model = sm.tsa.SARIMAX(data['Sales'], exog=scaled_exog_data, order=(0,0,0), trend='c', error_ar = 1)
results = model.fit()
results.summary()

<class 'statsmodels.iolib.summary.Summary'>
"""
                               SARIMAX Results
=====

Dep. Variable:                  Sales   No. Observations:      100
Model:                          SARIMAX   Log Likelihood:  -825.765
Date:                Mon, 10 Mar 2025   AIC:                  1663.530
Time:                      10:30:18   BIC:                  1679.161
Sample:                         0   HQIC:                  1669.856
                                  - 100

Covariance Type:                  opg
=====

              coef    std err        z     P>|z|    [0.025
0.975]
-----
intercept  2.022e+04    107.340   188.362      0.000      2e+04
2.04e+04
Base_Price -1017.6055    106.716    -9.536      0.000    -1226.764
-808.447
TV          322.4525     95.305     3.383      0.001     135.657
509.248
Radio       103.2478    116.714      0.885      0.376    -125.508
332.003

```

```

InStore      476.8354    106.885     4.461      0.000    267.344
686.327
sigma2       8.71e+05    1.17e+05     7.467      0.000    6.42e+05
1.1e+06
=====
=====
Ljung-Box (L1) (Q):           31.23   Jarque-Bera (JB):
3.10                           0.00   Prob(JB):
0.21
Heteroskedasticity (H):      1.77   Skew:
-0.33
Prob(H) (two-sided):         0.11   Kurtosis:
3.55
=====
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients
(complex-step).
"""

base_price = [15.64, 15.49]
radio = [279, 259]
instore = [41.50, 20.40]
tv = [103.44, 128.40]
testdf =
pd.DataFrame({'Base_Price':base_price,'TV':tv,'Radio':radio,'InStore':instore})
mean = data[['Base_Price','TV','Radio','InStore']].mean()
std = data[['Base_Price','TV','Radio','InStore']].std()
x = testdf[['Base_Price','Radio','InStore','TV']]
x_var = (x-mean)/std
x_var

{
"summary": {
  "name": "x_var",
  "rows": 2,
  "fields": [
    {
      "column": "Base_Price",
      "properties": {
        "dtype": "number",
        "std": 0.20250065839097342,
        "min": 0.39375199243717063,
        "max": 0.6801311699231664,
        "num_unique_values": 2,
        "samples": [
          0.39375199243717063,
          0.6801311699231664
        ],
        "semantic_type": "\",
        "description": "\n"
      }
    },
    {
      "column": "InStore",
      "properties": {
        "dtype": "number",
        "std": 1.091297933987541,
        "min": -0.8839343574329994,
        "max": 0.6593939814019196,
        "num_unique_values": 2,
        "samples": [
          -0.8839343574329994,
          0.6593939814019196
        ],
        "semantic_type": "\",
        "description": "\n"
      }
    },
    {
      "column": "Radio",
      "properties": {
        "dtype": "number",
        "std": 1.091297933987541,
        "min": -0.8839343574329994,
        "max": 0.6593939814019196,
        "num_unique_values": 2,
        "samples": [
          -0.8839343574329994,
          0.6593939814019196
        ],
        "semantic_type": "\",
        "description": "\n"
      }
    }
  ]
}

```

```
    "std": 0.1625632204768092, "min": 0.026553347336826096,  
    "max": 0.2564524584781773, "num_unique_values": 2,  
    "samples": [0.2564524584781773, 0.026553347336826096],  
    "semantic_type": "\\",  
    "description": "\"\\n      }\\n  },\\n  {\\n    \"column\":  
    \"TV\",\\n    \"properties\": {\\n      \"dtype\": \"number\",\\n      \"std\": 0.40542719580961206, \"min\": -0.8458698893485193,  
      \"max\": -0.27250925047967345, \"num_unique_values\": 2,  
      \"samples\": [-0.27250925047967345, 0.8458698893485193],  
      \"semantic_type\": \"\\\",  
      \"description\": \"\\n      }\\n    }\\n  ]\\n}\",  
    "type": "dataframe", "variable_name": "x_var"  
  
forecast = results.forecast(steps=2, exog=x_var)  
forecast  
100    19362.405485  
101    19405.838486  
Name: predicted_mean, dtype: float64
```